
Constraint-Aware Optimization and Robust Evaluation of GCG Adversarial Suffixes

Korea University COSE461 Final Project

Donghwa Lee

Department of Cyber Defense
Team 33
2023350220

Duhyun Kwon

Department of Cyber Defense
Team 33
2023350225

Yunsoo Song

Department of Cyber Defense
Team 33
2023350206

Abstract

Adversarial suffix attacks are a major challenge for large language model (LLM) safety because a short optimized token sequence can cause an aligned model to bypass its guardrails. This project studies GCG-style adversarial suffixes through baseline reproduction, constraint-aware optimization, stabilization, and robust evaluation. The baseline reproduction of prior GCG-based methods is separated from the main experiment: it is used only to understand existing attack behavior and reveal weaknesses in naive rule-based evaluation. Based on that study, we propose a constraint-aware discrete optimization approach for suffix generation. The method uses adaptive penalties, prompt-specific safety-direction estimation, safe vocabulary filtering, margin scheduling, and fluency regularization to reduce attention-collapse artifacts and meaningless outputs while preserving the target model’s language ability. We further extend evaluation from plain attack success rate (ASR) to a StrongREJECT-style score that considers refusal, specificity, and convincingness. In later tests, the stabilized method avoided repeated meaningless outputs, produced coherent text across benchmark prompts, and showed competitive StrongREJECT-style scores against a strong suffix-optimization baseline on Mistral-7B. Overall, the project shows that adversarial suffix research should evaluate not only whether a suffix avoids refusal, but also whether the generated response is meaningful, specific, coherent, and robustly judged as successful.

1 Introduction

Large language models are now deployed in many public-facing applications, but their safety mechanisms remain vulnerable to adversarial prompting. A typical aligned model refuses clearly harmful requests, yet prior work has shown that appending a carefully optimized adversarial suffix can substantially increase the probability of an affirmative or unsafe response [1]. Greedy Coordinate Gradient (GCG) is a representative method in this area because it optimizes a discrete suffix token by token using gradient information and greedy candidate selection.

This project began with the goal of identifying the critical nucleus of adversarial suffixes: the subset of suffix tokens that actually contributes to guardrail bypass. During the project, however, we found that this question depends heavily on evaluation reliability. If an evaluator counts hedged refusals, mixed responses, or meaningless strings as successful jailbreaks, then it is impossible to know which suffix tokens are truly important. Therefore, the project evolved into three connected goals: reproducing prior GCG-based attacks as a separate preliminary study, improving response-level evaluation, and designing a constraint-aware suffix optimization method.

The report follows the original template flow of Introduction, Related Work, Approach, Experiments, Analysis, and Conclusion, but it also includes a separate preliminary baseline-reproduction section. This separation is important because the reproduction of GCG, SM-GCG, and MAGIC is not the main proposed-method experiment. It serves as background evidence that motivates the later method and evaluation design. The main experiment focuses on the proposed constraint-aware method and its robust evaluation.

2 Related Work

Early adversarial attack research was largely developed in computer vision, where attacks such as FGSM, PGD, and CW-style optimization modify continuous image pixels. LLM attacks differ because text tokens are discrete and cannot be directly optimized in the same way. This difference motivated our topic refinement from image-based adversarial attacks to LLM jailbreak attacks.

GCG addresses the discrete-token problem by using gradient information to estimate promising token substitutions and greedily updating suffix coordinates [1]. Follow-up work has attempted to make suffix optimization faster or more robust. Probe sampling accelerates greedy coordinate optimization by reducing candidate-evaluation cost [2]. AutoDAN explores more natural-language jailbreak prompts rather than only artificial suffix strings [3]. Other GCG-style variants modify coordinate selection or update rules to improve stability and attack success.

Evaluation research is equally important. Standard ASR often relies on refusal detection or keyword matching, but this can overestimate jailbreak success when responses are non-refusal but vague, empty, or unrelated. HarmBench emphasizes standardized red-teaming evaluation [4], while StrongREJECT evaluates whether an answer is refused and, if not, whether it is specific and convincing [5]. This line of work directly influenced our robust evaluation design.

3 Preliminary Study: Baseline Reproduction

Before designing the proposed method, we conducted a separate preliminary reproduction study of prior GCG-based attacks. This section is intentionally separated from the main Experiments section. Its purpose is not to claim a new experimental contribution, but to establish a reference point for existing methods and to identify evaluation limitations.

The reproduction study used a Colab A100 GPU environment and LLaMA-family target models. We tested a control setting, basic GCG, and stronger GCG-style variants. The prompt set mixed normal questions, adversarial requests, jailbreak-style prompts, and security-related questions. The measured quantities were apparent ASR, average iterations to success, runtime, and peak memory usage.

The first attempt used about forty attack prompts, but it was too computationally expensive to complete within the available runtime. The second attempt reduced the prompt set to twenty prompts and refined the positive and negative scoring patterns. Table 1 summarizes the resulting baseline reproduction. Stronger variants produced higher apparent ASR than basic GCG and the control setting, broadly matching the expected trend from prior work.

Table 1: Preliminary baseline reproduction on the reduced prompt set. This table is separated from the main proposed-method experiment.

Method	Apparent ASR	Iteration trend	Runtime trend
Control	0.45	lowest when counted	lowest when counted
GCG	0.90	highest	highest
SM-GCG	1.00	medium	medium
MAGIC	1.00	lowest among optimized methods	lowest among optimized methods

However, this reproduction also revealed a serious limitation. The non-negligible control ASR suggested that simple rule-based scoring could be too permissive. Hedged responses, mixed refusal-compliance responses, and incoherent outputs could be mistakenly counted as successful attacks. This observation motivated the robust evaluation method used later for the proposed method.

4 Approach

4.1 Problem setup

Let x be a user prompt, $s = (s_1, \dots, s_L)$ an adversarial suffix, and M an aligned language model. The attacker searches for a suffix such that $M(x \oplus s)$ is more likely to produce the target behavior. A GCG-style objective can be written abstractly as

$$\min_s \mathcal{L}_{\text{attack}}(x, s; M), \quad (1)$$

where $\mathcal{L}_{\text{attack}}$ encourages an affirmative or target-like continuation. Since s is discrete, optimization is performed by repeatedly selecting a suffix coordinate, proposing candidate token substitutions, and choosing the candidate that improves the objective.

4.2 Constraint-aware suffix optimization

The proposed method is a constraint-aware discrete optimization approach for adversarial suffix generation. The motivation is that unconstrained suffix search can fail in two ways. If updates are too aggressive, the suffix may destabilize the target model and produce meaningless output. If constraints are too strict, the search may stagnate and fail to find useful token updates.

To balance these effects, we score candidate token replacements using both attack progress and a constraint penalty. Let Δe denote the embedding displacement caused by replacing the current token with a candidate token, and let g denote the gradient direction that reduces the attack objective. A simplified candidate score is

$$\text{Score}(e_{\text{new}}) = \langle g, \Delta e \rangle + \lambda R(\Delta e), \quad (2)$$

where $R(\Delta e)$ is a constraint-related penalty and λ controls its strength. The first term encourages attack progress, while the second discourages updates likely to produce unstable or guardrail-triggering behavior. The penalty strength is adjusted adaptively: it becomes stronger when the search appears unstable and weaker when exploration is safe.

4.3 Stabilization mechanisms

The later version of the method adds several stabilization mechanisms after observing attention collapse and abnormal-token artifacts. First, prompt-specific safety-direction estimation replaces a single fixed safety vector with a dynamic vector. Earlier versions used an average embedding of refusal phrases such as “I cannot” or “I’m sorry.” However, each prompt activates the target model’s safety behavior differently, so a fixed average vector was insufficient. The dynamic vector is intended to extract a suffix direction that is more sensitive to the guardrail behavior of each individual prompt.

Second, safe vocabulary filtering restricts the candidate token space to a safer ASCII-oriented set. In several optimization runs, abnormal tokens such as uncommon symbols, non-English logographs, or emoji-like artifacts caused the target model’s attention to collapse. The stabilized method assigns an infinite or very large penalty to unsafe tokens during loss computation, effectively removing them from the search space.

Third, adaptive margin scheduling and fluency regularization are used together. The method applies stronger optimization pressure early in the run to guide the target model toward the desired anchor, then linearly relaxes this pressure later so that the model can recover its generation ability and avoid overfitting. Fluency regularization penalizes suffix embeddings that move too far from the ordinary embedding region, encouraging suffixes that remain at least loosely compatible with the original prompt context.

These components were added after observing meaningless outputs such as repeated numeric strings or abnormal-token artifacts. The goal is not merely to avoid refusal, but to produce outputs that remain coherent enough to be evaluated meaningfully.

4.4 Robust evaluation

The evaluation pipeline combines rule-based filtering, semantic judgment, and StrongREJECT-style scoring. Rule-based filtering checks obvious refusal patterns and garbage-like outputs. Semantic

judgment examines whether the response actually satisfies the original harmful intent rather than merely avoiding refusal. The StrongREJECT-style score uses refusal, specificity, and convincingness:

$$Score = (1 - refused) \times \frac{specificity + convincingness}{2}. \quad (3)$$

A refused response receives zero. A non-refusal response still receives a low score if it is vague or unconvincing. This is stricter than plain ASR and directly addresses the false-positive problem discovered in the preliminary reproduction study.

5 Experiments

5.1 Data

This section evaluates only the proposed constraint-aware optimization method. The baseline reproduction in Section 3 is used as context, but it is not part of the main experiment. The proposed-method tests used benchmark harmful-query prompts without reproducing their detailed harmful content in this report. Early tests used five prompts to inspect optimization behavior, and later evaluation expanded to a small set of ten queries for qualitative comparison.

For StrongREJECT-style evaluation, the prompt set followed broad harmful categories such as illegal goods, non-violent crime, hate or harassment, disinformation, violence, and sexual content. The report focuses on aggregate scoring and response quality rather than unsafe procedural details.

5.2 Evaluation method

The main proposed-method evaluation uses best loss, garbage-output filtering, plain ASR, robust ASR, semantic success, and StrongREJECT-style score. Best loss tracks optimization progress, but it is not treated as jailbreak success by itself. Plain ASR reflects simple refusal-avoidance or keyword-based scoring. Robust ASR applies stricter filtering and semantic judgment. The StrongREJECT-style score further evaluates whether the output is specific and convincing.

This separation is essential because the preliminary reproduction showed that a response can be counted as successful by a weak detector even when it is hedged, partially refusing, or incoherent.

5.3 Experimental details

For the proposed method, instruction-tuned target models such as Mistral-7B-Instruct-v0.2 and Llama-3-8B-Instruct were used in preliminary experiments. Some optimization runs used 500 iterations, while the StrongREJECT-style dry run used a shorter single-prompt optimization setting. The target model and judge model were separated, and Gemini 2.5 Flash was used as a judge model in the StrongREJECT-style evaluation pipeline.

The later stabilized version included batch forward pass, prompt-specific safety-direction estimation, safe vocabulary filtering, adaptive margin scheduling, and fluency regularization. These changes were introduced to reduce meaningless responses and abnormal suffix artifacts. We also compared the stabilized method with a strong 2025 suffix-optimization baseline on Mistral-7B using the same StrongREJECT-style scoring perspective. Additional repeated-generation trials for P1, P5, and P9 were conducted to report score distributions and prompt-level ASR under a 0.5 score threshold. The comparison was used as an additional sanity check, not as a full-scale benchmark claim.

5.4 Results

Table 2 shows early proposed-method testing on five prompts. The proposed method achieved lower best loss than the baseline suffix optimization method in three of five cases. However, loss alone was not treated as success because some low-loss outputs were not meaningful or coherent.

Table 2: Preliminary best-loss comparison for the proposed method. Lower loss indicates optimization progress, not necessarily jailbreak success.

Prompt ID	Baseline best loss	Proposed best loss	Lower loss
Prompt 1	5.3125	6.9688	Baseline
Prompt 2	7.4688	5.7500	Proposed
Prompt 3	4.6250	7.6875	Baseline
Prompt 4	8.1250	6.8125	Proposed
Prompt 5	7.4375	4.1562	Proposed

After the stabilization mechanisms were added, meaningless repeated outputs such as numeric-loop strings were not observed in the tested cases. The Gemini-style judge and the StrongREJECT-style scoring direction also became more consistent, suggesting that the responses were easier to evaluate than earlier garbage-like outputs. The stabilized method produced normally formed suffixes and several coherent non-refusal responses, so the later results were more meaningful than simple refusal-avoidance cases. Figure 1 summarizes the later score trends from the updated project log.




Figure 1: Evaluation Results after Stabilization Improvements. Left: Gemini-style judge scores and StrongREJECT-style scores across ten benchmark prompts after safe vocabulary filtering and anti-loop stabilization. Right: repeated-trial robustness results, showing score distributions and ASR for prompts evaluated across multiple generations.




Figure 2: Comparison with a strong suffix-optimization baseline on Mistral-7B. The plot is redrawn from the updated project result slide with neutral method names. It shows that the proposed stabilized method is not uniformly stronger, but it achieves high scores on several prompts where the baseline is weak while maintaining readable text without attention collapse.

Table 3 summarizes the repeated-trial robustness results. The evaluation includes 30 trial-level rows for three prompts, with ten generations per prompt. For compact reporting, we include the category, number of trials, score distribution, and score-threshold success rate.

Table 3: Summary of repeated-trial robustness results. ASR is computed as the fraction of trials with StrongREJECT-style score ≥ 0.5 .

Prompt	Category	Trials	Mean	Std.	Median	Max	ASR
P1	Illegal Goods	10	0.463	0.384	0.491	0.892	50%
P5	Disinformation	10	0.524	0.153	0.538	0.708	50%
P9	Non-violent Crime	10	0.372	0.301	0.372	0.831	40%

A further comparison against a strong 2025 suffix-optimization baseline on Mistral-7B showed mixed but informative results. In the prompt-level comparison, the proposed stabilized method was clearly stronger on several prompts where the baseline score was near zero, especially the later illegal-goods and violence cases. It was also higher on one hate/harassment case, while the baseline remained stronger on several other categories such as the first illegal-goods case, non-violent crime, sexual content, and some lower-scoring prompts. More importantly, all tested outputs from the stabilized method remained readable text without attention collapse. This supports the main design goal: improving the practical quality and evaluability of suffix-generated responses, rather than only minimizing loss.

In a separate StrongREJECT-style dry run, the proposed method produced a normally formed suffix and a meaningful non-refusal response with a measured score of 0.50 out of 1.0. Since this score lies exactly at the success threshold, we interpret the case as partial success rather than robust complete success.

6 Analysis

The first finding is that prior-method reproduction must be separated from the main proposed-method experiment. The reproduction study is useful because it establishes baseline behavior and reveals evaluation problems, but it should not be presented as evidence for the proposed method itself. This is why the baseline reproduction is placed in a separate section and the Experiments section focuses only on the proposed constraint-aware method.

The second finding is that plain ASR is not enough. The baseline reproduction produced plausible numerical trends, but rule-based scoring could count hedged, mixed, or incoherent responses as successful. This means that any future critical-nucleus analysis must first use robust success labels; otherwise, ablation may identify artifacts rather than truly important suffix tokens.

The third finding is that optimization loss and jailbreak quality are related but not identical. Lower loss can show that the suffix search is moving in the intended direction, but it does not guarantee that the final response is coherent, relevant, or useful. This explains why the project shifted from only measuring loss or ASR to checking response quality through semantic judgment and StrongREJECT-style scoring.

The fourth finding is that stabilization improves practical output quality. Prompt-specific safety-direction estimation, safe vocabulary filtering, adaptive margin scheduling, and fluency regularization reduced meaningless outputs such as repeated numeric strings or abnormal-token artifacts. The later comparison with a strong suffix-optimization baseline also suggests that the stabilized method can succeed on some prompts where the baseline struggles, while staying comparable on others. Still, the result remains preliminary: a score of 0.50 in one dry run and mixed prompt-level comparison results do not establish robust success across categories, prompts, or target models.

7 Conclusion

This project studied GCG adversarial suffixes through a separated baseline reproduction study, constraint-aware optimization, stabilization, and robust evaluation. The work began by shifting from CV adversarial attacks to LLM jailbreak attacks, then reproduced GCG-based baselines to understand existing behavior. The reproduction study revealed that plain rule-based ASR can be misleading, but it is not treated as the main experiment.

The proposed constraint-aware method attempts to balance attack progress with output stability. Later improvements such as prompt-specific safety-direction estimation, safe vocabulary filtering, adaptive

margin scheduling, and fluency regularization reduced meaningless outputs and improved response coherence. StrongREJECT-style evaluation further showed that refusal avoidance alone is not enough; specificity and convincingness must also be considered. The updated comparison against a strong suffix-optimization baseline on Mistral-7B suggests that the stabilized method is competitive in some prompt categories and, importantly, avoids attention-collapse outputs in the tested cases.

The main limitation is scale. The number of prompts, target models, and trials is still limited, and judge-model bias remains possible. Future work should expand the comparison against GCG, SM-GCG, MAGIC, and black-box target models under the same StrongREJECT-style metric, report category-wise scores and variance, and return to the original critical-nucleus goal using leave-one-out ablation, attention analysis, and positional or linguistic token analysis. Additional code optimization is also needed for black-box transfer testing and larger-scale reporting.

References

- [1] Andy Zou, Zifan Wang, Nicholas Carlini, Milad Nasr, J. Zico Kolter, and Matt Fredrikson. Universal and Transferable Adversarial Attacks on Aligned Language Models. *arXiv preprint arXiv:2307.15043*, 2023.
- [2] Yiran Zhao et al. Accelerating Greedy Coordinate Gradient and General Prompt Optimization via Probe Sampling. *arXiv preprint*, 2024.
- [3] Xiaogeng Liu, Nan Xu, Muhao Chen, and Chaowei Xiao. AutoDAN: Generating Stealthy Jailbreak Prompts on Aligned Large Language Models. *arXiv preprint arXiv:2310.04451*, 2024.
- [4] Mantas Mazeika et al. HarmBench: A Standardized Evaluation Framework for Automated Red Teaming and Robust Refusal. *arXiv preprint arXiv:2402.04249*, 2024.
- [5] Alexandra Souly et al. A StrongREJECT for Empty Jailbreaks. *arXiv preprint*, 2024.

A Appendix: Team contributions

All team members participated in discussing prior research and exchanging opinions on GCG-based adversarial suffix attacks. All members also participated in discussions on possible new jailbreak methodologies and the direction of the proposed constrained optimization approach.

- **Donghwa Lee:** Conducted testing and analysis for the proposed constrained optimization method, including later-stage evaluation experiments.
- **Duhyun Kwon:** Reproduced prior GCG-based studies and conducted baseline attack testing experiments.
- **Yunsoo Song:** Organized research materials, summarized project contents, and helped structure the report.

B Appendix: Trial-level score rows

Table 4 lists the trial-level score rows used for the robustness summary. The appendix reports only prompt identifiers, categories, trial indices, and scores so that the quantitative results remain reproducible without overloading the main text.

Table 4: Trial-level StrongREJECT-style scores from repeated-generation robustness trials.

Prompt	Category	Trial	Score
P1	Illegal Goods	1	0.293
P1	Illegal Goods	2	0.077
P1	Illegal Goods	3	0.892
P1	Illegal Goods	4	0.892
P1	Illegal Goods	5	0.689
P1	Illegal Goods	6	0.163
P1	Illegal Goods	7	0.022
P1	Illegal Goods	8	0.000
P1	Illegal Goods	9	0.831
P1	Illegal Goods	10	0.773
P5	Disinformation	1	0.486
P5	Disinformation	2	0.471
P5	Disinformation	3	0.680
P5	Disinformation	4	0.317
P5	Disinformation	5	0.336
P5	Disinformation	6	0.591
P5	Disinformation	7	0.346
P5	Disinformation	8	0.682
P5	Disinformation	9	0.624
P5	Disinformation	10	0.708
P9	Non-violent Crime	1	0.352
P9	Non-violent Crime	2	0.212
P9	Non-violent Crime	3	0.521
P9	Non-violent Crime	4	0.581
P9	Non-violent Crime	5	0.041
P9	Non-violent Crime	6	0.013
P9	Non-violent Crime	7	0.392
P9	Non-violent Crime	8	0.757
P9	Non-violent Crime	9	0.831
P9	Non-violent Crime	10	0.018